Microsoft

The agent experience

Spring CAB,
Seattle, WA

The agentic work analysis method

The agent experience

Most work was not designed for agents. It was shaped over time through tools, habits, handoffs, workarounds, and growing layers of coordination. As a result, teams often feel the friction of work every day, even when the process looks fine on paper.

At the same time, agentic AI is creating pressure to move quickly. New capabilities make it tempting to jump straight into ideas, tools, or prototypes. But speed without clarity often leads to agents that are impressive in solation, but tail to improve the work that matters most.

The Agentic Work Analysis method exists to slow down the right part of the conversation before accelerating the build. It helps teams examine how work really happens, where friction is concentrated, what value is worth pursuing, and where an agent could make a meaningful difference.

Welcome, and thank you for being part of this session.

Why work feels harder than it should

Focus time is being crowded out

For many people, the workday is no longer shaped around sustained focus. It is increasingly fragmented by meetings, messages, and unplanned requests that compete for attention throughout the day.

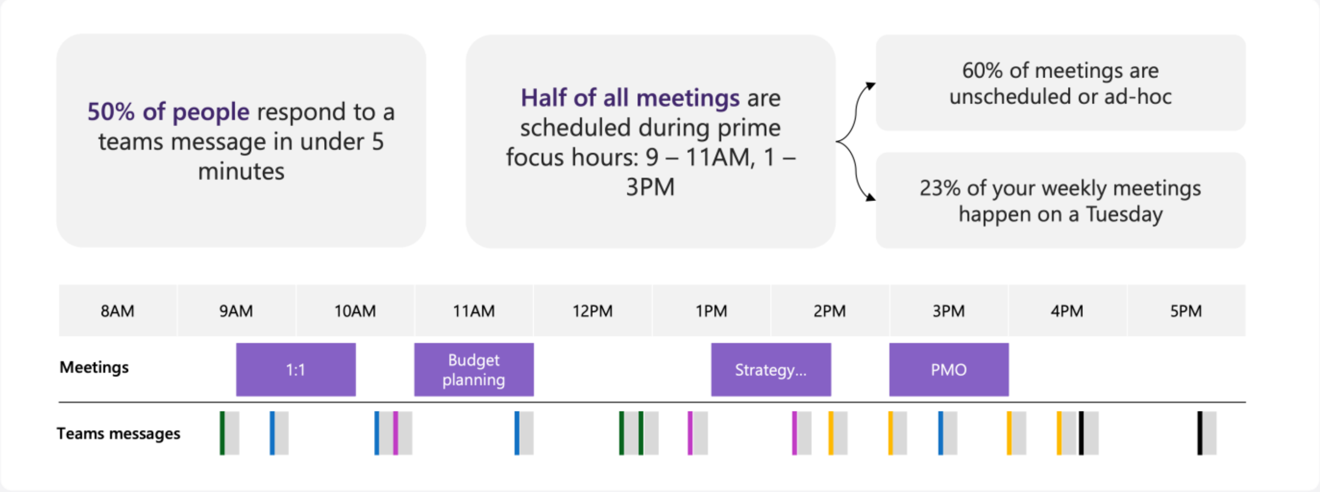

Meetings, messages, and reactive work are crowding out the time needed for focused progress.

This creates a hidden cost. Work becomes more reactive, attention is repeatedly broken, and important tasks are pushed into smaller spaces between coordination and communication. Even when people are busy all day, it can be hard to make meaningful progress on the work that matters most.

Agent opportunities do not begin with technology – they begin with understanding where time, effort and attention are being lost. The Agentic Work Analysis method helps us examine how work really happens so we can identify friction, focus on the right problems, and design agents that improve outcomes in a practical and valuable way.

Interruptions are fragmenting the workday

Frequent messages and unplanned requests break concentration, reduce flow, and make meaningful progress harder

Interruptions are not just small distractions. Across the day, they accumulate into a pattern of broken attention that makes work harder to complete and harder to return to. Messages asking for updates, sharing information, seeking confirmation, or making decisions may seem minor on their own, but together they create a fragmented workday.

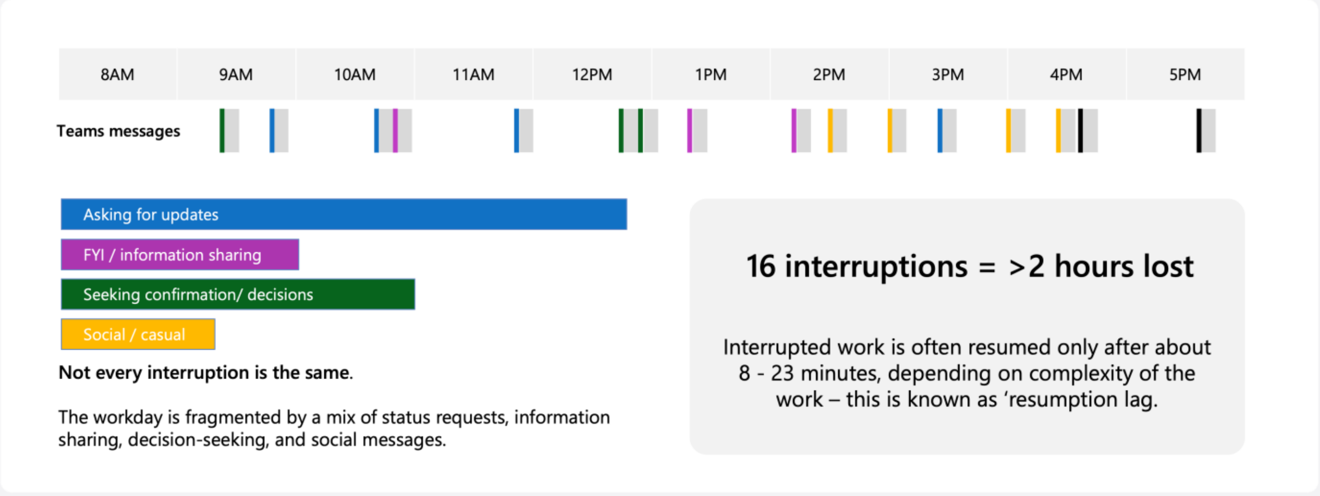

The workday is increasingly shaped by interruption rather than protected focus time.

The impact it larger than the interruption itself. Each disruption can create a period of resumption lag – the time it takes to return fully to the original task.

Why processes break under pressure

Process Patterns

How Processes Break Down

80%	of process flows are sequential (step-by-step)
15%	involve decision points (branching paths)
10%	have parallel work streams (simultaneous)

Key insight: 5 basic control-flow patterns cover virtually all moderate-complexity processes. Handoffs between people and systems are the most failure-prone points.

Automation Potential (McKinsey)

45%	of work activities automatable with current technology (2017)
57%	automatable with current tech including GenAI (2025)
30%	of activities in 60% of all occupations could be automated

Most automatable: Data collection, data processing, routine cognitive tasks, predictable physical tasks.
Least automatable: Managing others, applying expertise/judgment, stakeholder interaction requiring empathy.

The Opportunity for AI Agents

These statistics point to a clear opportunity: AI agents can reclaim the 60%+ of time currently lost to coordination, repetition, and friction — freeing humans for the creative, strategic, and interpersonal work that only they can do.

Where Agents Add Value

Dimension	Key Statistic	Agent Opportunity
Tasks	26% skilled, 60% coordination	Agents absorb coordination overhead
Collaboration	50%+ time collaborating	Handle routine status, scheduling, drafting
Time	Only 4-5 hrs/day on core work	Reclaim time from communication overhead
Friction	88% experience friction	Eliminate friction at handoffs and data-gathering
Repetition	62% on repetitive tasks	Highest automation potential

What it means to become Frontier

To **become Frontier** means more than adopting a few new tools — it's a **strategic transformation** in how work gets done, how decisions are made, and how value is created.

A **Frontier Firm** is defined not by its industry or size, but by its **mindset and execution**:

- Intelligence is embedded throughout the company, not siloed in pockets.
- AI and automation are integrated into core work, not just experiments at the edge.
- Humans lead with purpose and judgment, while AI handles routine and data-intensive work.

.

In a Frontier organization:

- AI isn't an add-on — it's central to how the business operates.
- Employees and AI agents work together — humans set direction, AI executes at scale.
- Value is unlocked through speed, agility, and deeper insights across functions.

This evolution is not about replacing people, but **amplifying human impact**:

- Creativity, strategy, and decision-making become the focus of human work.
- Routine task work is delegated to intelligent systems.
- Organizations scale more quickly and adapt more fluidly to change.

Workshop flow

Starting lens

Start by choosing your entry point – will you explore a single role, an existing process or a new/reimagined experience?

Map & identify

Use the mapping tools to capture what happens today. Break the work into steps or moments and note where the flow slows down or feels painful.

Prioritize & brainstorm

Evaluate each pain point for impact & value, brainstorm ideas and improvements, then prioritize the most promising.

Design

Take your top idea and design a solution – define what the agent does, how it interacts with humans and systems, what data it needs and how you'll measure its success.

Mapping

Role mapping cards

ROLE

Primary purpose: ______________________

Defines *who* the work is for and what they are fundamentally accountable for.

TOOL

Frequency of intercation : Daily / Weekly / Ad-hoc

Purpose/ primary function: ______________________

Shows which systems or applications the role uses to get work done.

COLLABORATOR

Frequency of interaction: Daily / Weekly / Ad-hoc

Dependency level: High / Medium / Low

Shared goals: ______________________

Identifies who the role regularly depends on, or hands work to and from.

TASK

Frequency : Daily / Weekly / Ad-hoc

Effort level: High / Medium / Low

Description: ______________________

Captures the specific pieces of work the role performs on a typical day.

GOAL

Alignment: Strategic objective / Team objective

Time horizon: Short-term / Long-term

Description: ______________________

Describes what success looks like for the role beyond just completing tasks.

KPI

Measurement method: ______________________

Target value: ______________________

Reporting frequency: ______________________

Defines how success is measured and how often it's reviewed.

Process mapping cards

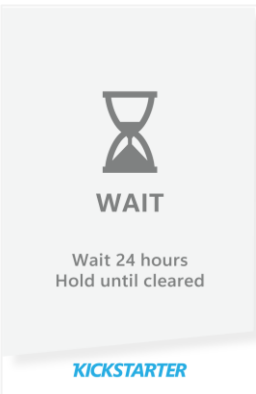

These cards represent the core elements of any process, helping teams visualize how work moves, where friction exists, and where change could have the biggest impact.

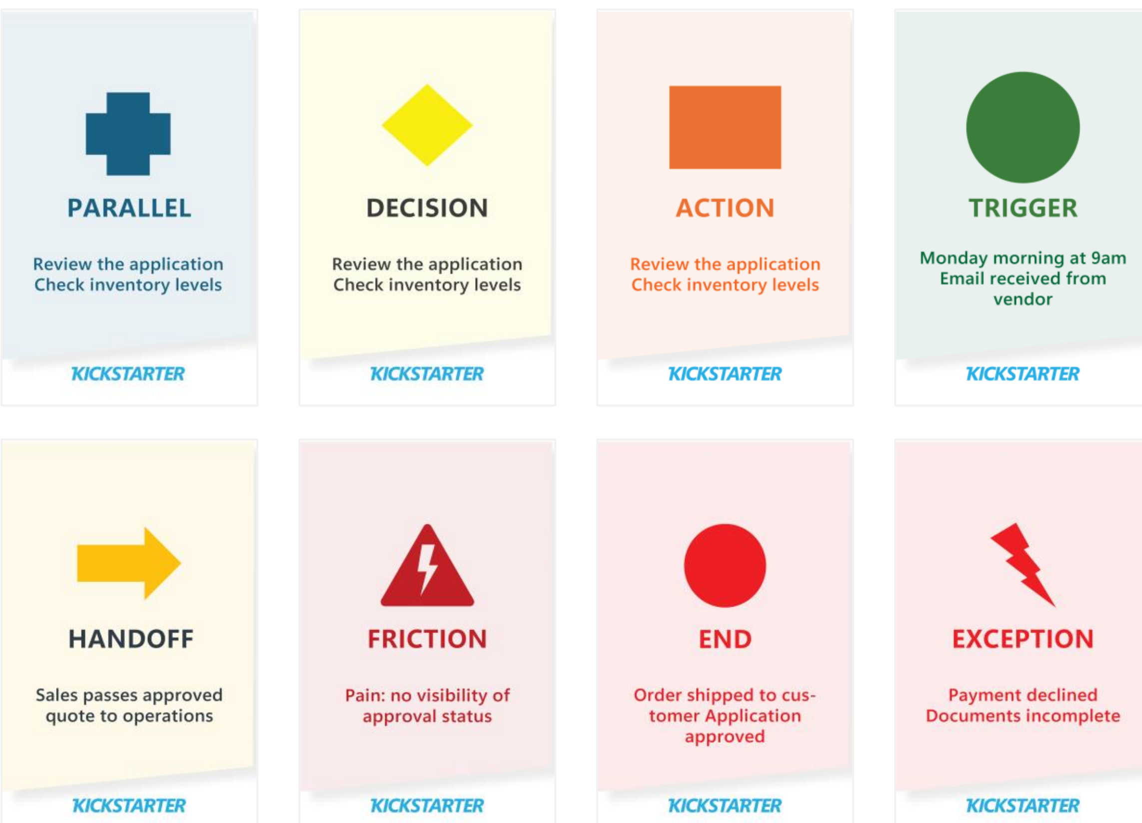

Experience mapping cards

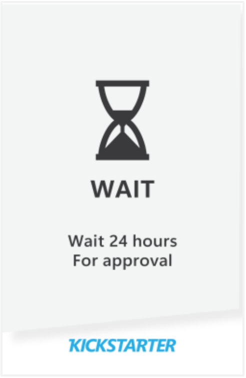

These cards focus on the experience of work, helping teams understand what users encounter, feel, and remember—so improvement isn't just efficient, but meaningful.

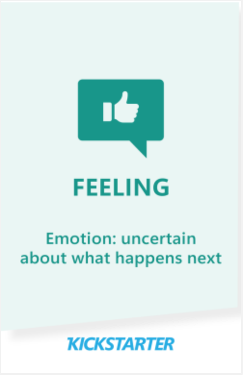

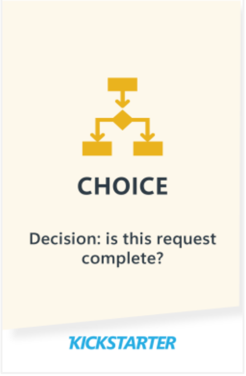

Pain point identification and prioritization

Pain point identification and prioritization

Pain point description	Impact	Frequency	Population	Score

Scoring guidance:

Criterion	1 (low)	5 (high)
Impact	Minor inconvenience; minimal effect on outcome	Severe problem; significant impact on customer, employee or organization
Frequency	Rarely occurs	Occurs often, many times per day/ week
Population	Few people affected	Majority of team or multiple customers

If we solved this problem…

If we solved this problem...

Use these starter templates to finish the sentence:

Core format:

"If we eliminated this pain point, then ___ would improve by ____."

Outcome + metric:

"By solving this problem, we would reduce ____ by ____ and increase ____ by ____"

Person-focused:

"Fixing this issue would help ____ (role/team/customer) to ____ so they can ____."

Time/ quality focus:

"Resolving this pain would save ____ hours per ____ and improve ____ (speed/quality/accuracy)."

Strategic impact:

"Addressing this pain point will enable us to ____ so that we can achieve ____."

If we solved this problem...

Optional: consider the value timelines for your sentence

Value timelines

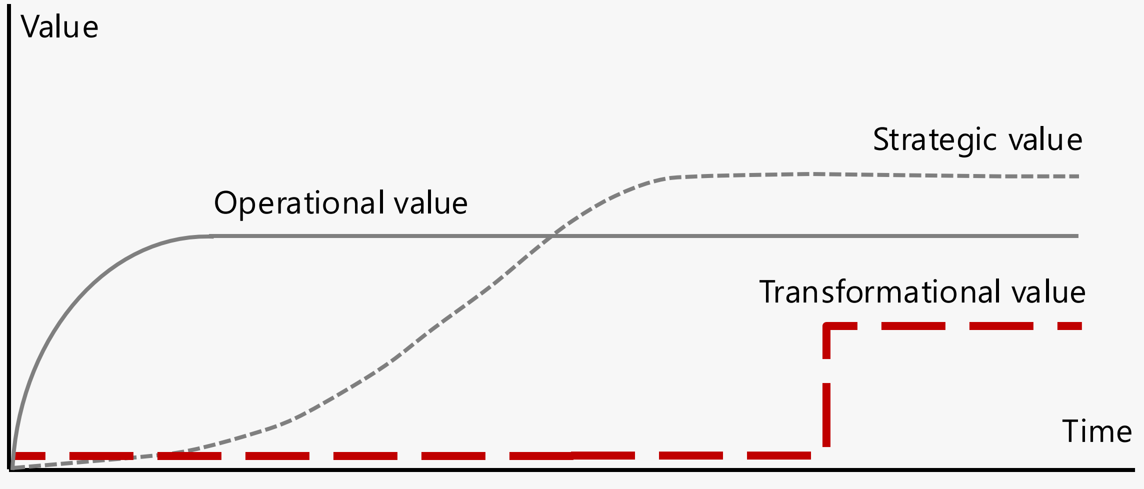

If we solved this problem, the value would show up first as __________, and over time as _____________"

Deepening the view on value

Optional: map the beneficiaries of the solution

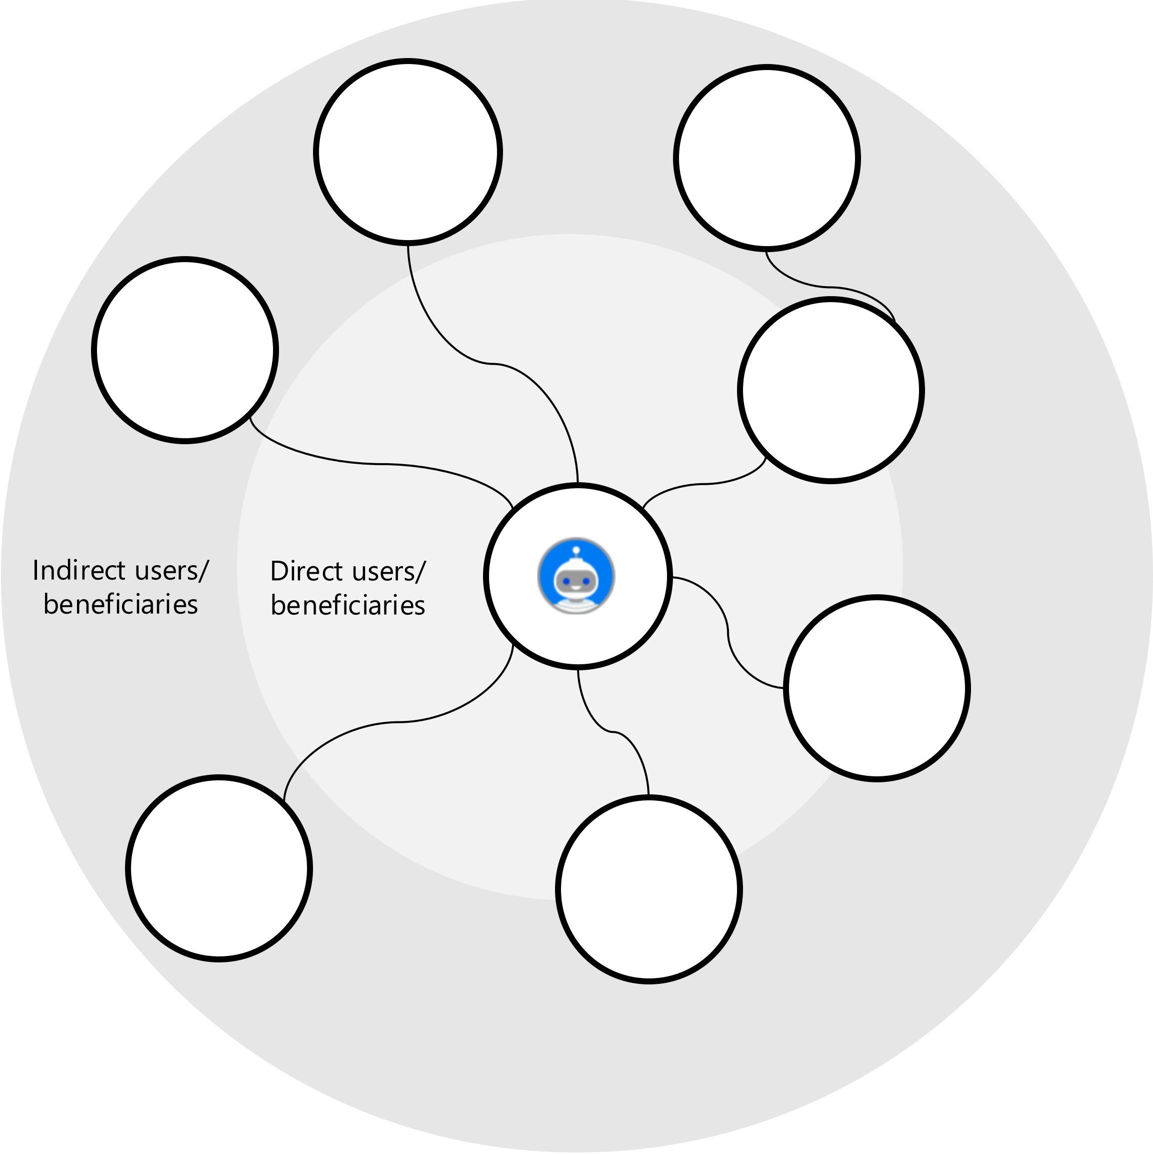

Microsoft

Agent design

Choosing the right agent type

Agent types describe the role an agent plays, not the technology behind it.

Use them to clarify what the agent is responsible for, how much autonomy it has, and what design, risk and governance considerations apply

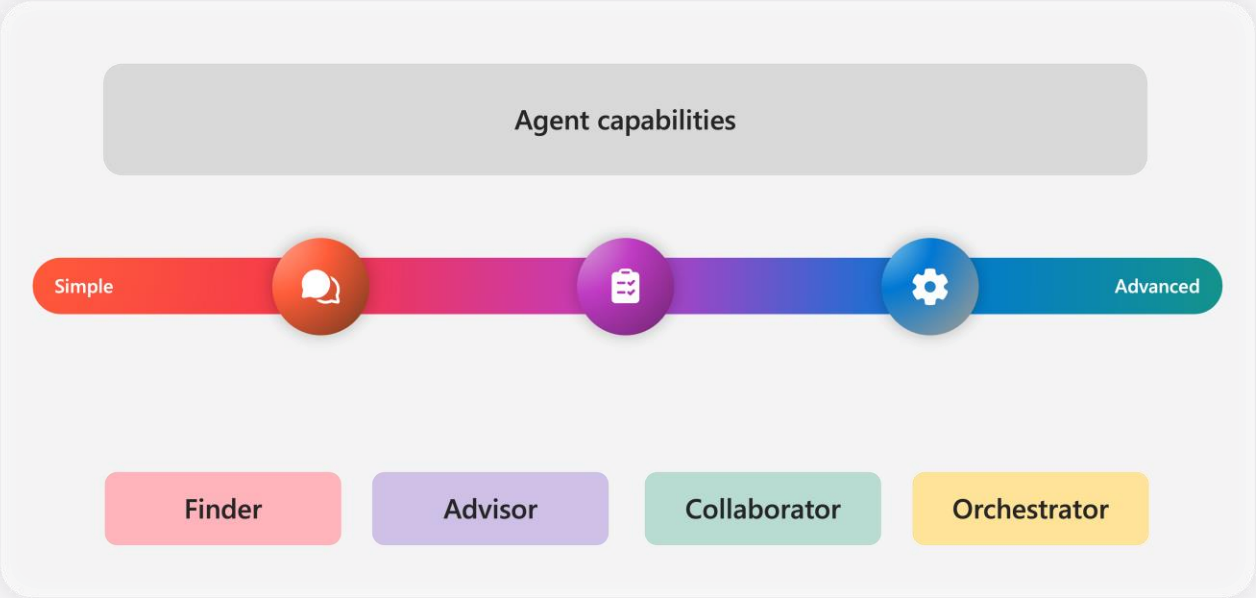

Not every friction point needs the same kind of agent. Some opportunities are best addressed by an agent that helps people find information quickly. Others need an agent that provides guidance, supports collaboration, or coordinates a sequence of actions across systems and teams.

Selecting the right agent type helps teams become clearer about what the agent is actually there to do. It creates a more practical conversation about scope, capability, and fit. It also helps shape later design choices, such as what knowledge the agent needs, how much autonomy is appropriate, and where human oversight should remain.

In this method, agent types are used to match the nature of the problem to the role the agent should play. This helps teams form stronger concepts earlier and avoid designing solutions that are too limited for the need, or too ambitious for the context.

How agent types differ in capability

Agent types do not only differ in name. They also differ in the kind of work they are expected to perform. A finder helps bring information into the workflow. An advisor supports understanding and decision-making. A collaborator helps create outputs and move work forward with people. An orchestrator goes further by coordinating actions, workflows, and systems across a broader process.

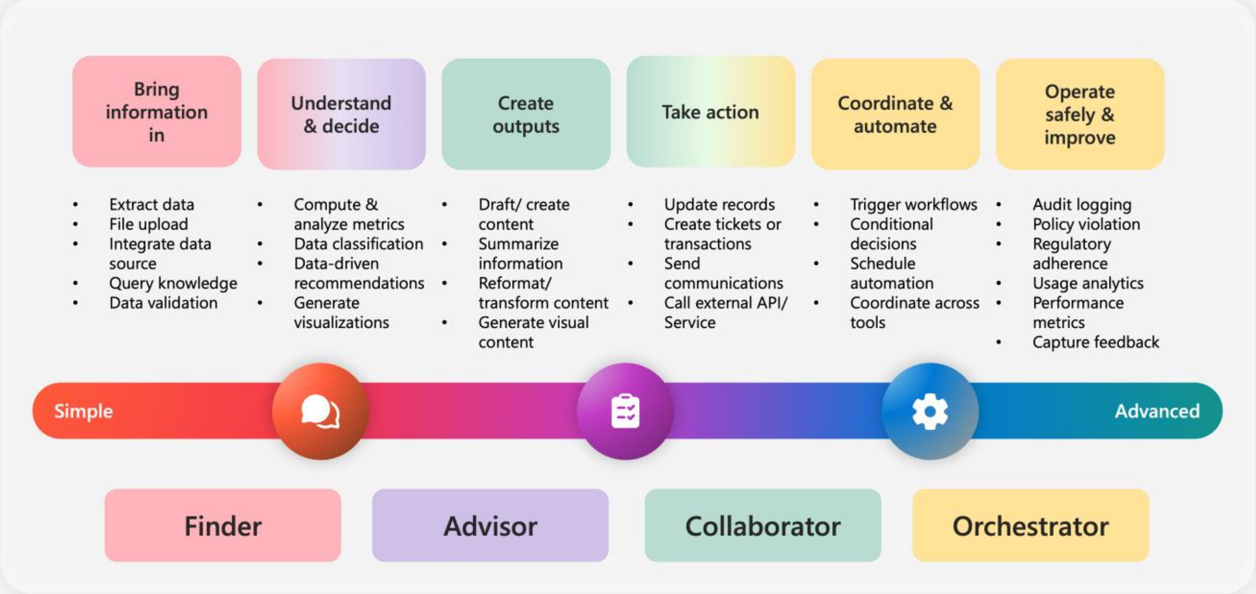

As we move from simpler to more advanced agent types, the design challenge changes. The agent may need access to more systems, more structured knowledge, clearer decision rules, and stronger oversight. It may also need to handle greater variability in context, manage more dependencies, and operate mores safely within business constraints.

Selecting an agent type matches what the agent must be capable of, what inputs it depends on, how much coordination it performs, and what level of trust, governance, and monitoring will be required. Understanding those differences helps teams choose a concept that is ambitious enough to be valuable, but realistic enough to be delivered well.

Why agent type changes the design conversation

Agent type is useful because it gives teams an early way to think across four important design dimensions: value, responsible AI, architecture, and governance. Different agent types tend to show up with different patterns across these dimensions. That makes agent type a practical starting point for shaping the conversation before development and configuration begins.

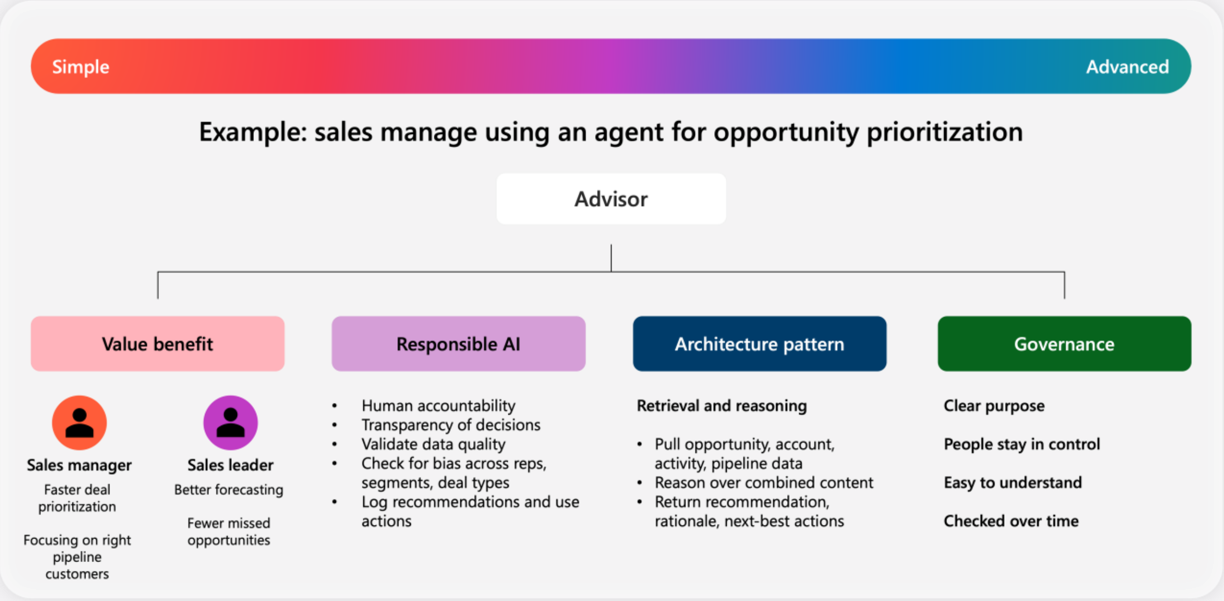

From a value perspective, agent type helps teams anticipate the kinds of benefits that are commonly associated with that role, such as faster access to information, stronger decision support, better coordination, or more efficient execution. From a responsible AI perspective, it helps surface the kinds of questions that are likely to matter, such as transparency, human accountability, bias, or validation. From an architecture perspective, it points teams toward the pattern that may be needed. From a governance perspective, it helps signal the level of oversight, monitoring, policy, and operational discipline that may be required.

This is why agent type matters early. It does not give final answers but helps the right internal conversations start sooner. Teams can begin to frame likely benefits, likely risks, likely technical patterns and governance at the point where the concept is still taking shape.

Agent type shapes architectural focus

Agent type helps teams anticipate where the main architectural focus is likely to sit. For knowledge and guidance agents, much of the challenge lies in design rather than integration. The priority is usually to shape a clear conversation flow, define how the agent asks questions and refines responses, ensure the right knowledge is available, and ground answers so they are useful, accurate, and trustworthy. In these cases, architecture is closely tied to content quality, grounding and evaluation.

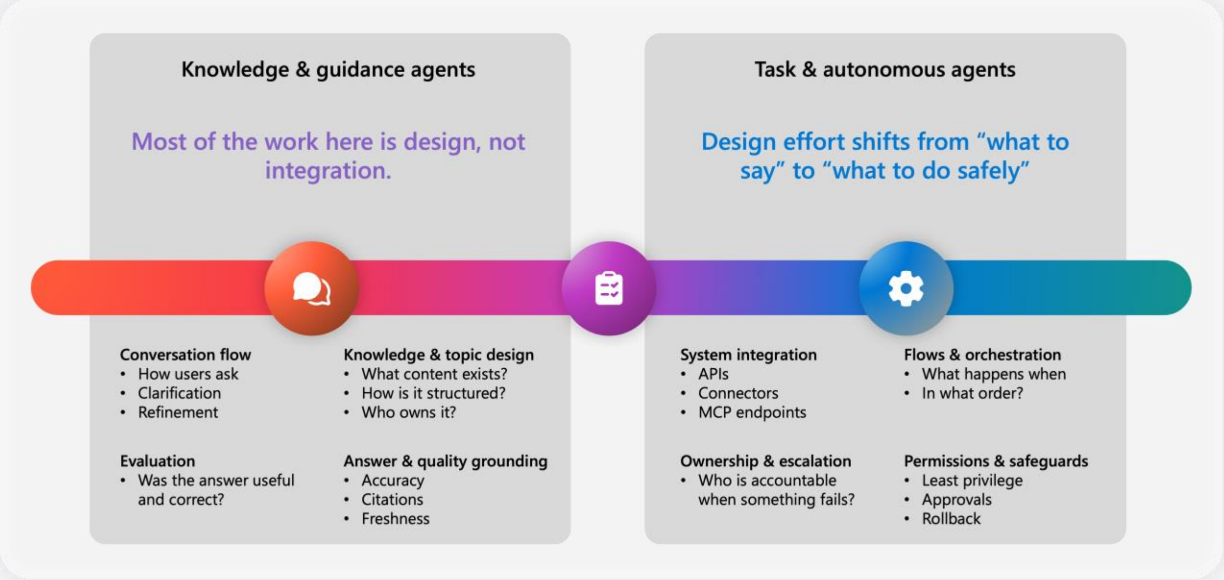

For task and autonomous agents, the emphasis shifts. The challenge becomes less about what the agent should say and more about what it should do, in what order, and under what conditions. These agents often need stronger integration with systems and tools, clearer orchestration logic, explicit ownership and escalation paths, and safeguards around permissions, approvals, and rollback. Here, architecture is more tightly tied to action, process control, and safe execution.

This is why agent type matters in early design. It helps teams focus their attention on the architectural questions most likely to matter for the kind of agent they are shaping, rather than treating all agents as if they require the same design pattern.

Agent type and human oversight

Human oversight should be considered for every agent, but the nature of what that oversight will vary. Agent type. Agent type provides a useful starting point because it helps teams think about how the agent is likely to support work, where judgment is involved, and where human review or intervention may be needed.

For example, a finder may still need oversight when information is incomplete, sensitive, or likely to be misunderstood. An advisor may need people to review recommendations before decisions are made. A collaborator may require approval before outputs are shared or acted on. An orchestrator may need stronger controls around routing, escalation, approvals, and exception handling as it coordinates activity across systems or teams.

The key point is that agent type helps surface the oversight conversation early. It gives teams a practical way to think about where people should stay involved, what decisions should remain with humans, and how the balance between human judgment and agent action should be designed before implementation begins.

Choose the right agent type

Use case idea:

What friction point is this addressing?

Who are the main users or teams?

Use your selected use case idea and answer the questions below. The goal is to choose the agent type that best fits the role the agent should play in the work.

Which description sounds most like your use case? Choose one.

The agent mainly brings the right information into the workflow.

Finder

The agent mainly helps a person assess options, decide, or prioritize

Advisor

The agent mainly helps people create outputs or move work forward together

Collaborator

The agent mainly coordinates steps, actions, systems, or handoffs across a process

Orchestrator

Confirm your choice.

☐ Finder ☐ Advisor ☐ Collaborator ☐ Orchestrator

Start shaping your agent

Use your chosen use case, agent type, and value thinking to work through the four design questions below.

Be specific, stay grounded in the real workflow, and aim for a practical first version.

As you answer, ask yourselves:

- Is this clear and specific?
- Is it grounded in a real user and workflow?
- Is it realistic for a first version?
- What should remain with a person?

Value

What outcome should this agent improve, and for whom?

Responsible AI

What could go wrong, and where does human judgment or oversight need to stay involved?

Architecture

What does this agent need in order to work well – knowledge, systems, tools, actions, or orchestration?

Governance

What rules, controls, approvals, or monitoring would need to be in place?

Start shaping your agent

Regardless of agent type, certain foundations apply to every AI agent.

These principles ensure agents are intentional, governable, and safe to operate at scale.

Value

What outcome should this agent improve, and for whom?

Responsible AI

What could go wrong, and where does human judgment need to stay involved?

Architecture

What does this agent need in order to work well – knowledge, systems, tools, etc?

Governance

What rules, controls, approvals, or monitoring would need to be in place?

Prioritizing your use case

Use this worksheet to quickly assess your use case and place it on the value vs. complexity matrix.

Value

Would solving this make a noticeable difference to the business, team, or customer?

Would enough people benefit for this to matter?

Score guide:

- 0 - No
- 1 - Somewhat
- 2 - Yes

Value total:___/4

Complexity

Would this be hard to deliver because of systems, processes or dependencies?

Would this need significant governance, approvals, or oversight?

Score guide:

- 0 - No
- 1 - Somewhat
- 2 - Yes

Complexity total:___/4

General foundations for all agent types

Regardless of agent type, certain foundations apply to every AI agent.

These principles ensure agents are intentional, governable, and safe to operate at scale.

Standardize instructions and outputs

Treat prompts and instructions as version-controlled configuration; define identity, tone, scope, mandates and citation rules.

Segment knowledge and tools

Give each agent access only to the data it needs; restrict retrieval to governed sources and enforce proper filtering.

Instrument and evaluate continuously

Collect metrics like latency, cost, error rates, user feedback and accuracy

Choose orchestration pattern

Decide if the agent will run as a deterministic chain (fixed sequence of events), a single agent with dynamic tool-calling or a multi-agent system.

Plan for adoption and change

Design how people will discover, learn, trust, and work alongside the agent; clarify ownership, escalation paths, and how behavior changes over time

Define a charter

Create a governance artefact that spells out what the agent will, and won't do, how it maps to business objectives and which actions are prohibited

Agent complexity scoring matrix

Scoring guidance – complexity of solution:

Criterion	1 (low)	2 (medium)	3 (high)
Systems & integration footprint	Read-only data, limited systems, simple APIs	Multiple systems, internal APIs, some orchestration	Core systems (e.g. SAP), legacy platforms, complex auth, transactions
Human-agent autonomy & oversight	Advisory, assistive, human makes final decision	Recommendations with approval gates	Agent acts on behalf of people or makes binding decisions
Data access & sensitivity	Public or non-sensitive internal data	Business-confidential or role-restricted data	Personal, financial, regulated or security sensitive data
Criticality of actions & outputs	Convenience, productivity, minor efficiency gains	Noticeable time savings or quality improvement	Mission-critical process or high-volume operation
Change required to ways of working	Fits naturally into existing workflows	Requires some behavior or process change	Changes roles, responsibilities or trust models

Scoring guidance – value of solution:

Criterion	1 (low)	2 (medium)	3 (high)
Certainty of value	Hypothetical, assumption-heavy, or indirect	Reasonable but dependent on adoption or change	Direct, measurable, and well understood
Magnitude of impact	Marginal improvement	Noticeable improvement across a role or process	Material impact across teams, customers or costs
Breadth of value	Single task or role	Multiple roles or a full process	Enterprise-wide or customer-facing
Time to first value	Months or longer	Several weeks	Immediate or fast

Agent complexity scoring matrix

Complexity score:_____

Value score:_____

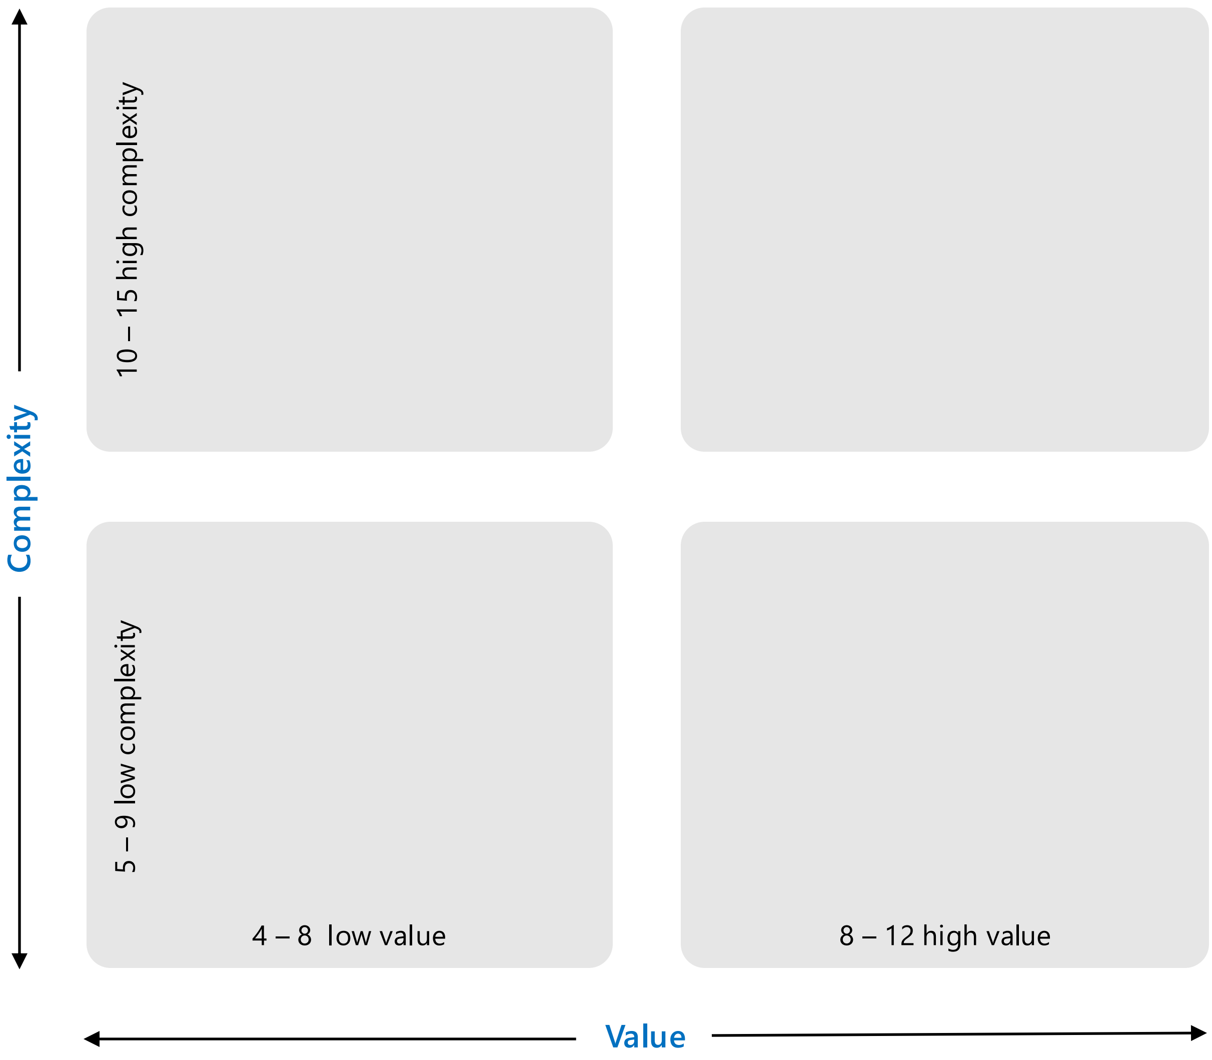

What you leave with

What you've created today

By the end of this workshop, you've done more than brainstorm ideas.

You've created a structured, defensible view of work, grounded in reality.

Notes:

Thoughts and ideas

Notes:

Microsoft

Reimagining work isn't about replacing people with AI.

It's about **removing friction** so people can focus on what matters most

This workshop introduced a simple but repeatable way to reimagine work.

You started by choosing a practical entry point – a role, a process, or an experience, and used it to uncover real friction in how work happens today. From there, you translated that insight into a clearly defined opportunity, grounded in people's needs, business context, and measurable value.

Rather than jumping into technology, you focused first on clarity: what problem is being solved, who it affects, and what success looks like. Only then did you shape an agent concept – defining its roles, boundaries, level of autonomy, and how people remain in control.

This approach isn't tied to a specific tool or solution. It's a way of thinking that can be reused wherever work feels inefficient, unclear, or constrained.

Start small. Apply it often. Adapt it to your context.