Microsoft

CAT

COPILOT ACCELERATION TEAM

CAT AI WEBINARS

Building an AI-ready organization

Business Value of AI Agents ROI, Impact, and Real-World Outcomes

Meet your host and speakers

Darian Nastvogel

Program Manager
CAT Scale Team

Niyonika Chhabra

PRODUCT MANAGER,
Copilot Agent Development

Moran Herman

SENIOR PRODUCT MANAGER,
Copilot Studio

Microsoft

Measuring what Matters: ROI and Business Value of AI Agents

ROI, Impact and Measurement Across the Agent Lifecycle | Microsoft Webinar | April 15, 2026

Agenda

1. **Why measurement is the missing link in enterprise AI**
   Why measurement is the missing link in enterprise AI

2. **The Business Value Framework**
   4 value pillars, ROI formula, and the recapture gap

3. **Top Use Cases & Measurement Blueprints**
   The 6 most common deployments and how to track them

4. **Microsoft Measurement Ecosystem**
   Built-in tools: Copilot Studio Analytics, Power CAT Kit, App Insights

The investment is real

AI agents are moving from experimentation to enterprise deployment, at scale, right now.

$10.9B
Global AI Agents market in 2026
Grand View Research

40%
of enterprise apps will embed agents by 2026
Gartner

$2.5T
Worldwide AI spending in 2026
Gartner

The question is no longer whether to build AI agents.

Everyone Is Building AI Agents

But Who’s Proving Value?

Three stages of the journey

Each section speaks to a different stage, from starting out to scaling across the organization.

01

PART ONE

Before Building

Beginning your agentic journey

"How do I build the right value foundation before I write a single line of code?"

02

PART TWO

Measuring Impact

Agent is live or nearly there

"What should I be measuring, and how do I connect it to business impact?"

03

PART THREE

Telling the Story

Data ready for stakeholders

"How do I turn metrics into a conversation my stakeholders find meaningful?"

1

Part One

Before building the agent

The value conversation starts before development does

Before a single line of code is written, how will you know if your agent delivered real value?

A QUESTION WORTH REFLECTING ON

Define value before you define the build

Every strong agent deployment begins with four questions answered in sequence, before any technical work begins.

Discovery → Envision → Build → Deploy → Measure & Optimize

01 Strategic Goal

What business challenge are we addressing?

02 Stakeholder Map

Who benefits, and what matters to them?

03 Success Metrics

How will we recognize that it worked?

04 Build

Now the team builds with a shared definition of value

When these four steps are addressed upfront, teams enter development with a clear picture of what success looks like and the baseline they need to prove it.

Choosing the right use case

Mapping use cases by impact and effort helps teams focus energy where value is most likely

	Low Effort	High Effort
High Impact	**Quick Wins** High impact, fast to build	**Strategic Investments** High impact, significant effort
Low Impact	**Nice to Have** Low impact, easy to build	**Reconsider** Low impact, high effort

SUGGESTED KPIs TO DEFINE EARLY	
Deflection Rate	Target: ____%
Resolution Rate	Target: ____%
CSAT Score	Target: ____/5
Monthly Savings	Target: $____
Escalation Rate	Target: <____%
Adoption Rate	Target: ____%

Agreeing on these targets early gives the team a shared foundation for measuring impact.

2

Part Two

During and after building the agent

With an agent live, the focus shifts to capturing what it delivers

With the agent deployed and usage growing, how does a team connect that activity to business outcomes?

A PATTERN WORTH RECOGNIZING

What most teams find challenging about measurement

THREE AREAS TO STRENGTHEN

01

Measurement that continues past launch

Teams often measure during pilots then step back once live. Treating ROI as ongoing makes the value story stronger.

Embed measurement into the deployment workflow for continuous tracking.

02

Connecting activity to business outcomes

Sessions and user counts show usage. Connecting those to cost, quality, or revenue is where the case becomes compelling.

Each KPI should be traceable to a value pillar the business cares about.

03

Translating time saved into value created

Time saved creates value when it flows into higher-value work. 50% productivity recapture rate is a widely used baseline.

Start with the Forrester 50% baseline and refine over time.

The Business Value Framework

Where business value comes from, and how to connect each pillar to something measurable.

Discovery → Envision → Build → Deploy → Measure & Optimize

01

Efficiency Value

Time saved × recapture rate × labor cost

The most accessible pillar. A 50% recapture rate keeps it grounded and defensible.

02

Quality Value

Error reduction + compliance consistency + rework avoidance

Often the largest driver in regulated industries. Benefits from early instrumentation.

03

Revenue Value

Conversion lift + retention improvement + new revenue enabled

Highest ceiling. Connecting agent telemetry to CRM and sales data is key.

04

Strategic Value

Decision velocity + employee confidence + competitive positioning

Increasingly quantifiable through NPS, satisfaction signals, and adoption trends.

The ROI formula

A single framework for the team

$$\text{AGENT ROI} = \frac{[\text{Sum of Pillar Benefits} \times \text{Adoption Rate} \times \text{Recapture Rate}] - \text{Total Cost of Ownership}}{\text{Total Cost of Ownership}} \times 100$$

Adoption Rate

Share of target users actively engaging, not just enrolled

Total Cost of Ownership

Licensing, development, integration, change management, maintenance

Recapture Rate

50% Forrester baseline: saved time that converts to productive output

Sum of Value Pillars

Efficiency, Quality, Revenue, and Strategic value quantified

Top 6 Use Cases and the Metrics that Matter

?

WHERE TEAMS ARE DEPLOYING

Where Organizations Are Deploying Agents Today

Top 6 Use Cases

Use Case	*Primary KPIs*	*Baseline Before Launch*	*Pillars*
HR & Employee Ops	**Deflection rate, resolution time, manager hrs saved, ticket volume**	**Avg handling time, weekly ticket volume, escalation rate, HR headcount ratio**	**Efficiency · Quality**
Customer Service	**Deflection %, first-contact resolution, CSAT/NPS delta, avg handle time**	**Cost per interaction, daily contact volume, avg handle time, escalation rate**	**Efficiency · Revenue**
Finance & Invoice	**Processing time, straight-through %, error rate, exception volume**	**Time per invoice, error rate, FTE hrs per 100 invoices, exception queue depth**	**Efficiency · Quality**
Procurement & Intake	**Intake cycle time, touchless PO %, supplier wait time, approval SLA**	**Intake cycle time, FTE hrs per request, approval turnaround, rework rate**	**Efficiency · Quality**
Knowledge & Docs	**Hours saved monthly, accuracy rate, retrieval time, repeat queries**	**Time to find info, escalation rate for knowledge gaps, content freshness**	**Efficiency · Strategic**
IT & Dev Productivity	**Ticket deflection %, dev productivity lift, MTTR, code-gen acceptance**	**Tickets per engineer/week, MTTR, dev hours on L1 issues, sprint velocity**	**Efficiency · Strategic**

3

Part Three

Telling the story

Metrics in a dashboard become a business case when they reach the right audience

Microsoft Measurement Ecosystem

Three Tiers of Measurement Maturity

TIER 1

Copilot Studio Analytics

Built-in. Zero config. Session data, resolution rates, deflection, AI-grouped themes from day one.

TIER 2

Copilot Studio Kit

App Insights, Power BI dashboards, extended history, agent value summaries across 10 benefit categories

TIER 3

Azure App Insights

Custom telemetry for latency, errors, business events. KQL queries and alerts for production deployments.

A closer look at Copilot Studio Analytics

Copilot Studio Analytics Goal

Help you understand how your Copilot agents **perform** and provide **guidance** for continuous improvement.

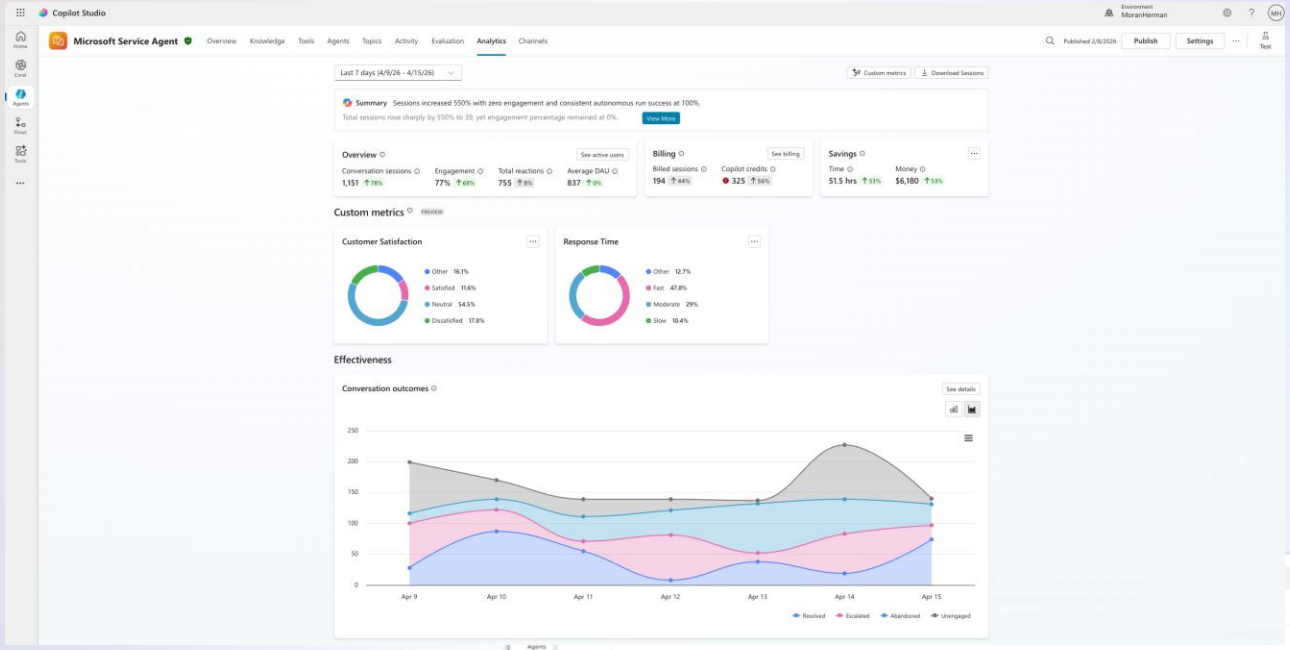

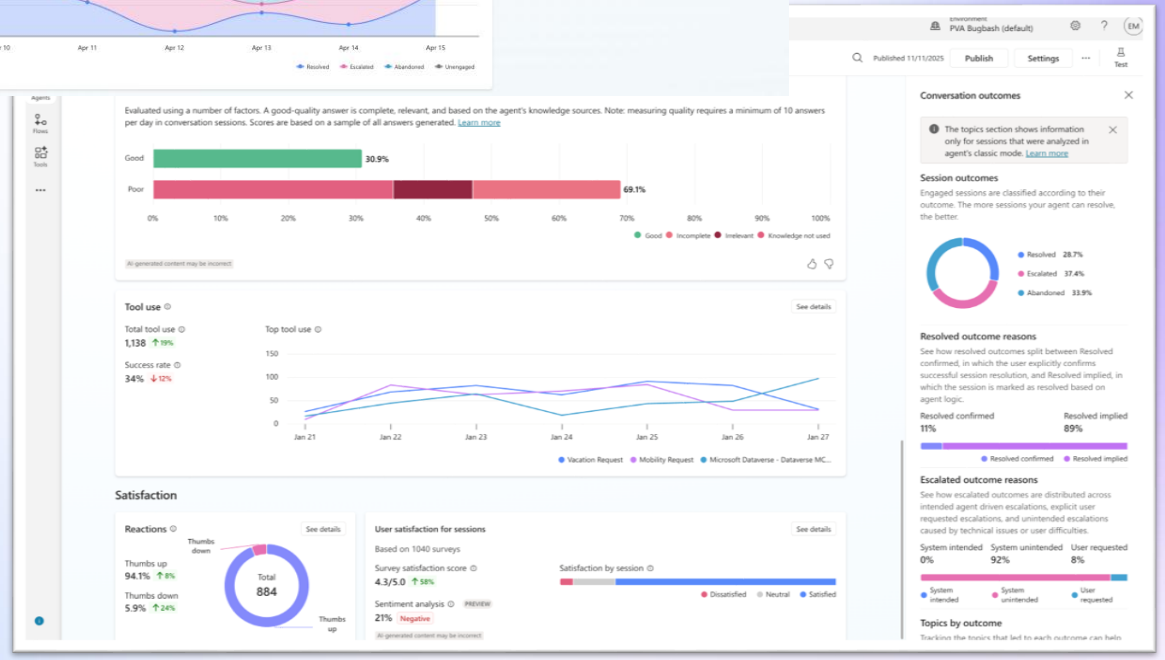

Copilot Studio

Environment
EranManor

EM

Home
Agents
Flows
Tools

Employee Self Service

Overview | Knowledge | Tools | Agents | Topics | Activity | Evaluation | **Analytics** | Channels

Published 1/29/2026 | Publish | Settings | Test

Last 7 days (1/25/26 - 1/31/26)

Download Sessions

Summary PREVIEW Sessions increased 550% with zero engagement and consistent autonomous run success at 100%.

Total sessions rose sharply by 550% to 39, yet engagement percentage remained at 0%. View More

Overview

Conversation sessions	Engagement	Satisfaction score	Average DAU	Copilot Credits used
1,507 ↑21%	72% ↑56%	4.3/5.0 ↑95%	179 ↑7%	307 ↑15%

Savings

Calculate savings

Start tracking savings to measure the time and money that your agent saves.

Effectiveness

Conversation outcomes

See details

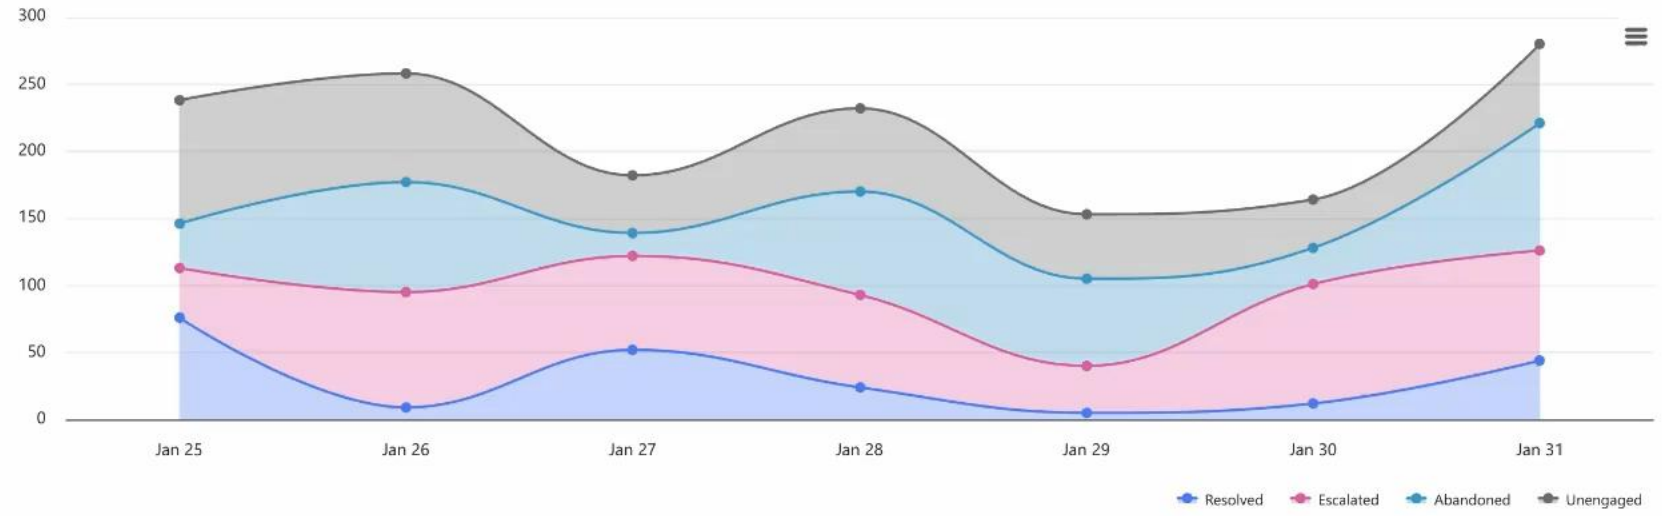

Use

Themes PREVIEW

See all (14)

Analyze your users' questions by auto-classified themes, and see how well your agent responds. Learn more Preview terms

Analytics Savings

Define and track the **time** and **money** saved by your agent to measure and communicate the ROI to stakeholders.

Savings can be defined per successful agent **run** or per **tool** execution.

<u>For example</u>:

Each successful run of an IT self-service agent saves an avg. 20 minutes of human representative time.

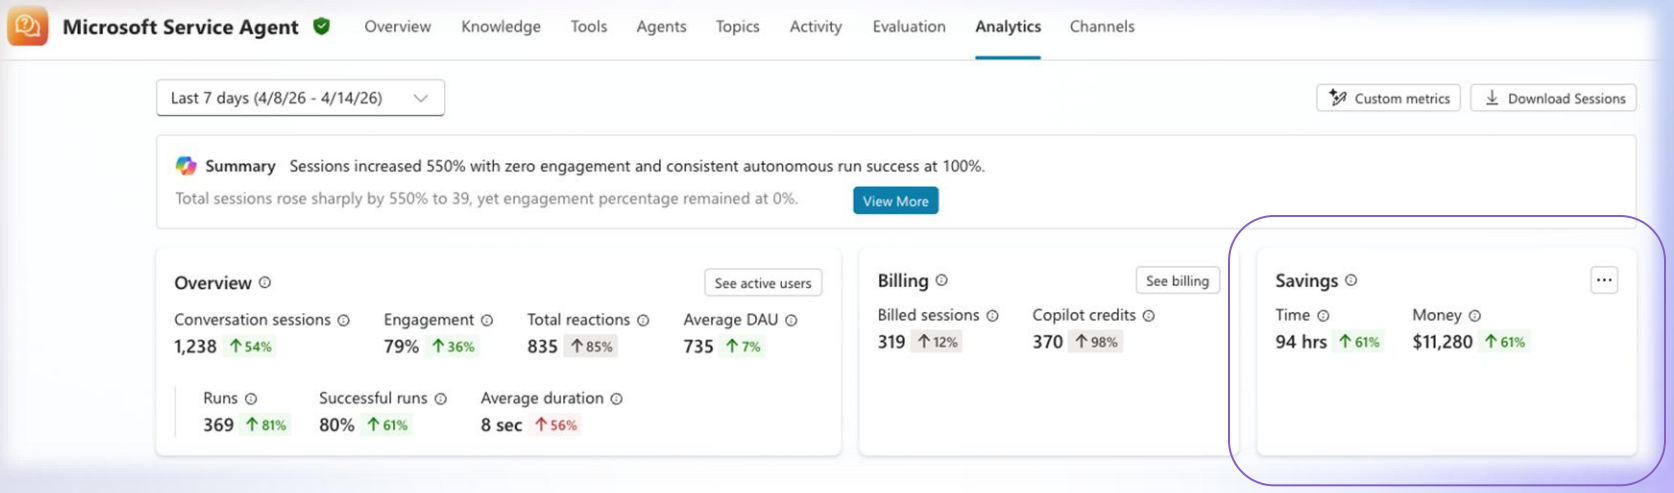

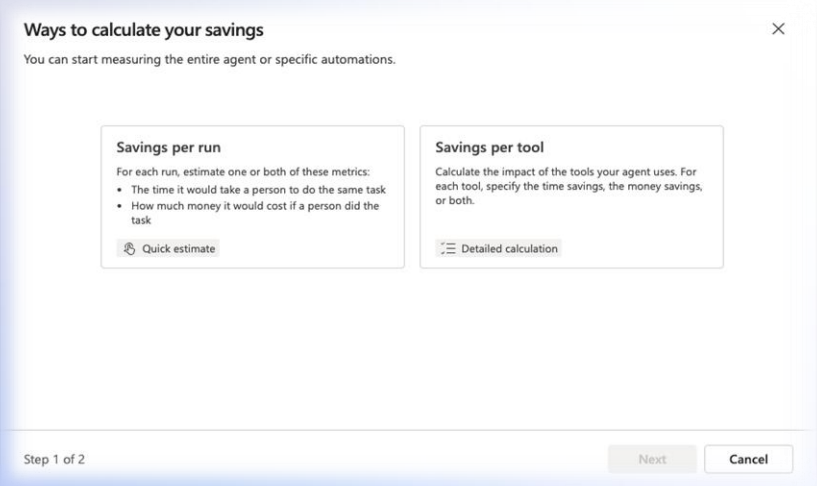

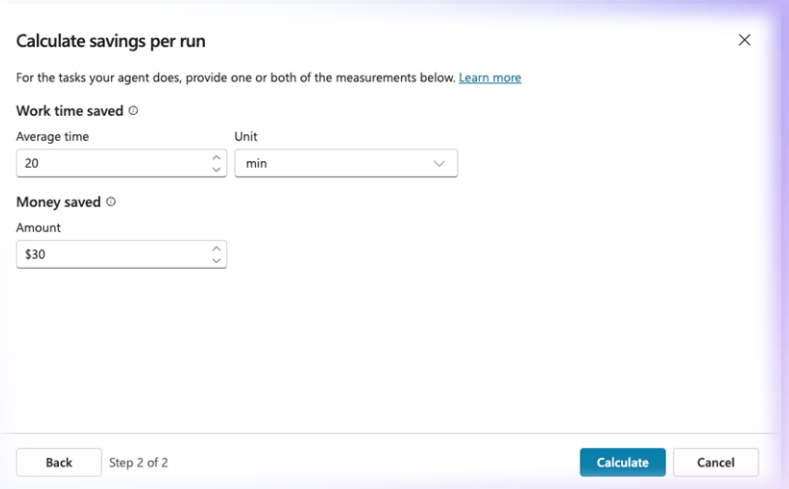

Custom Metrics

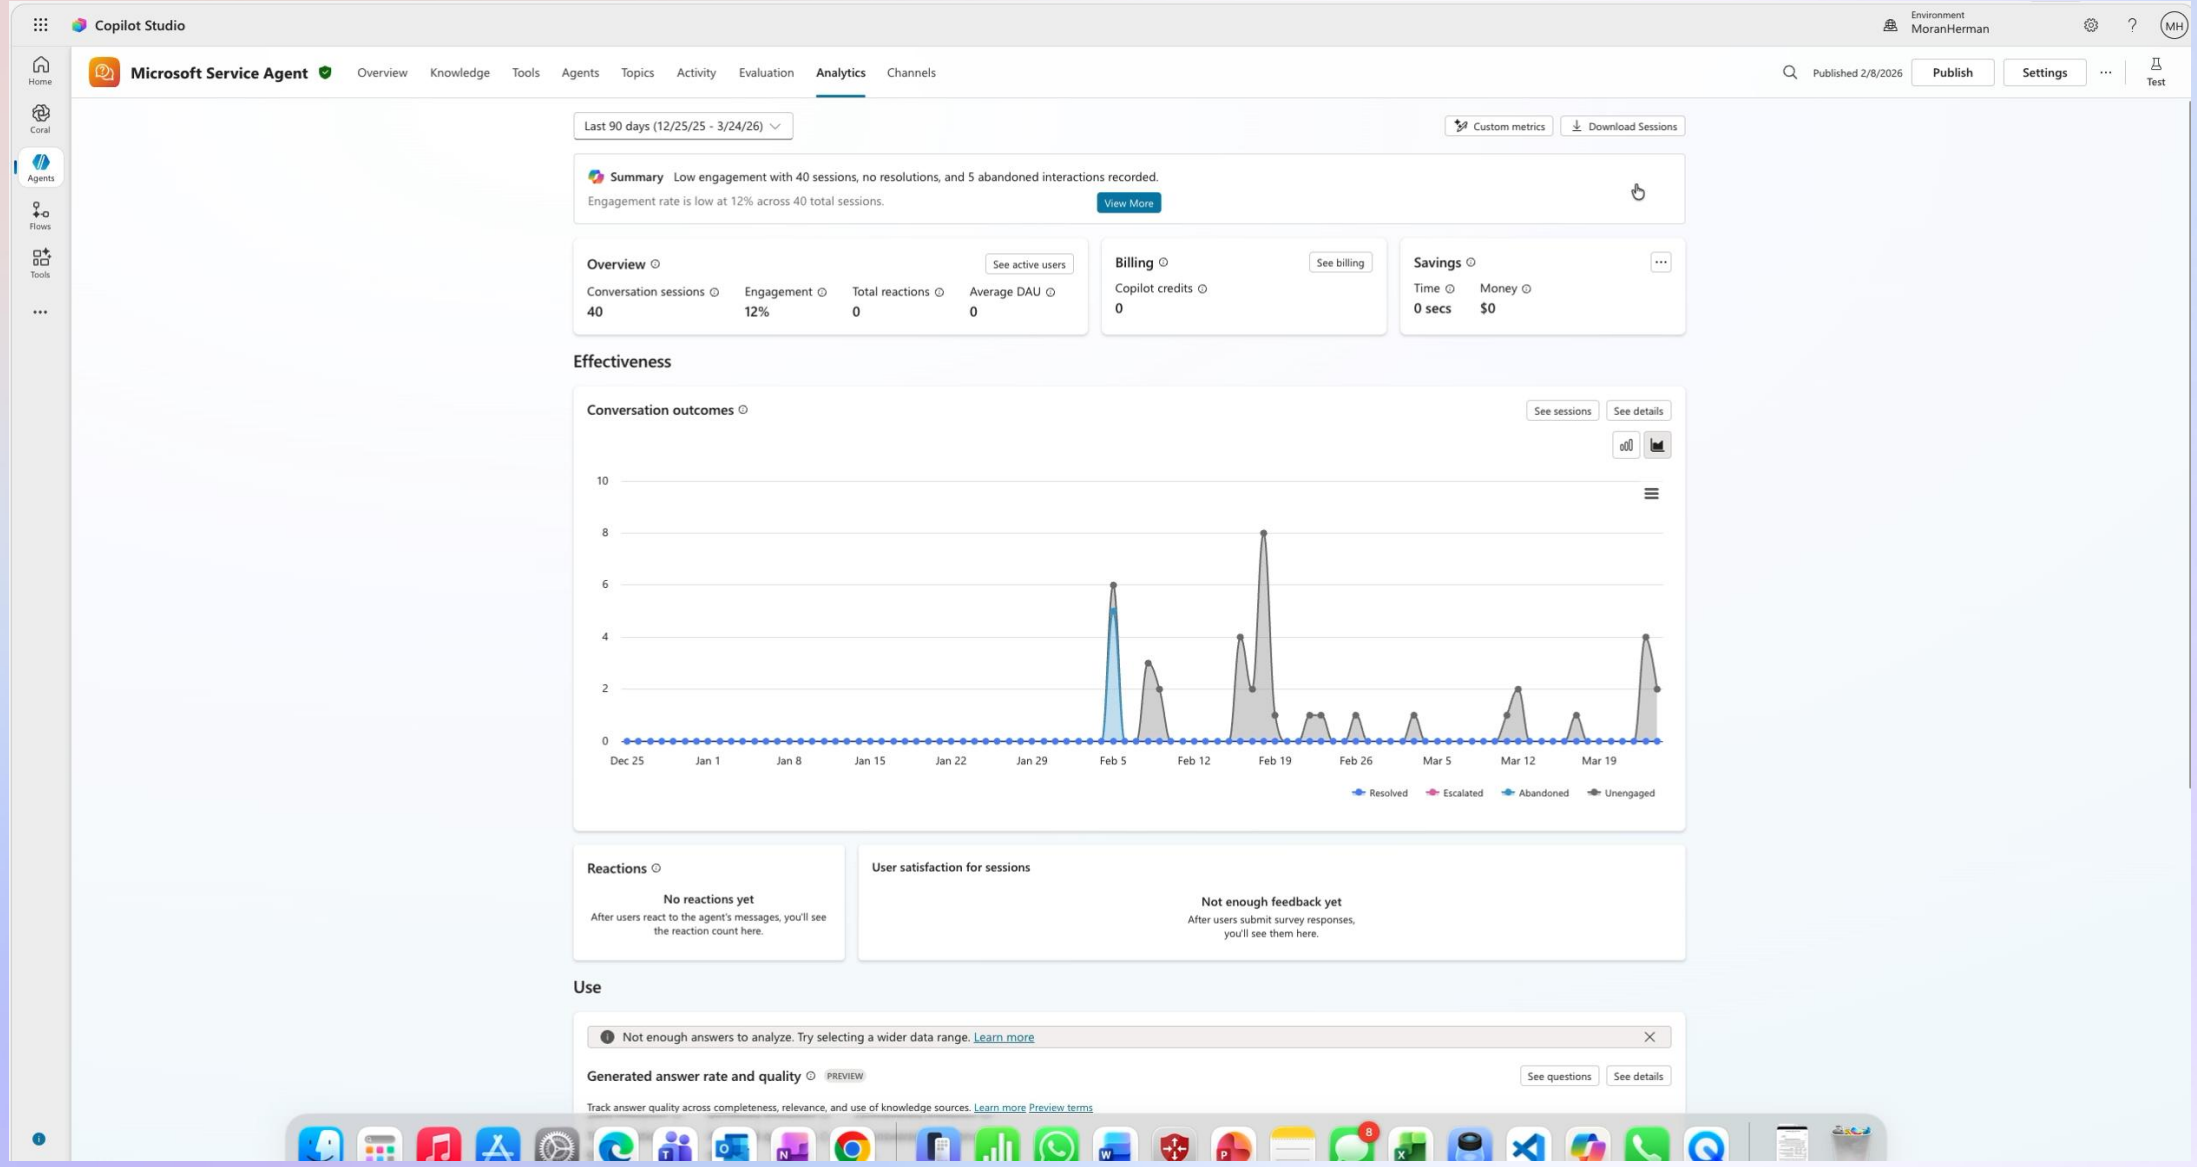

Custom Metrics

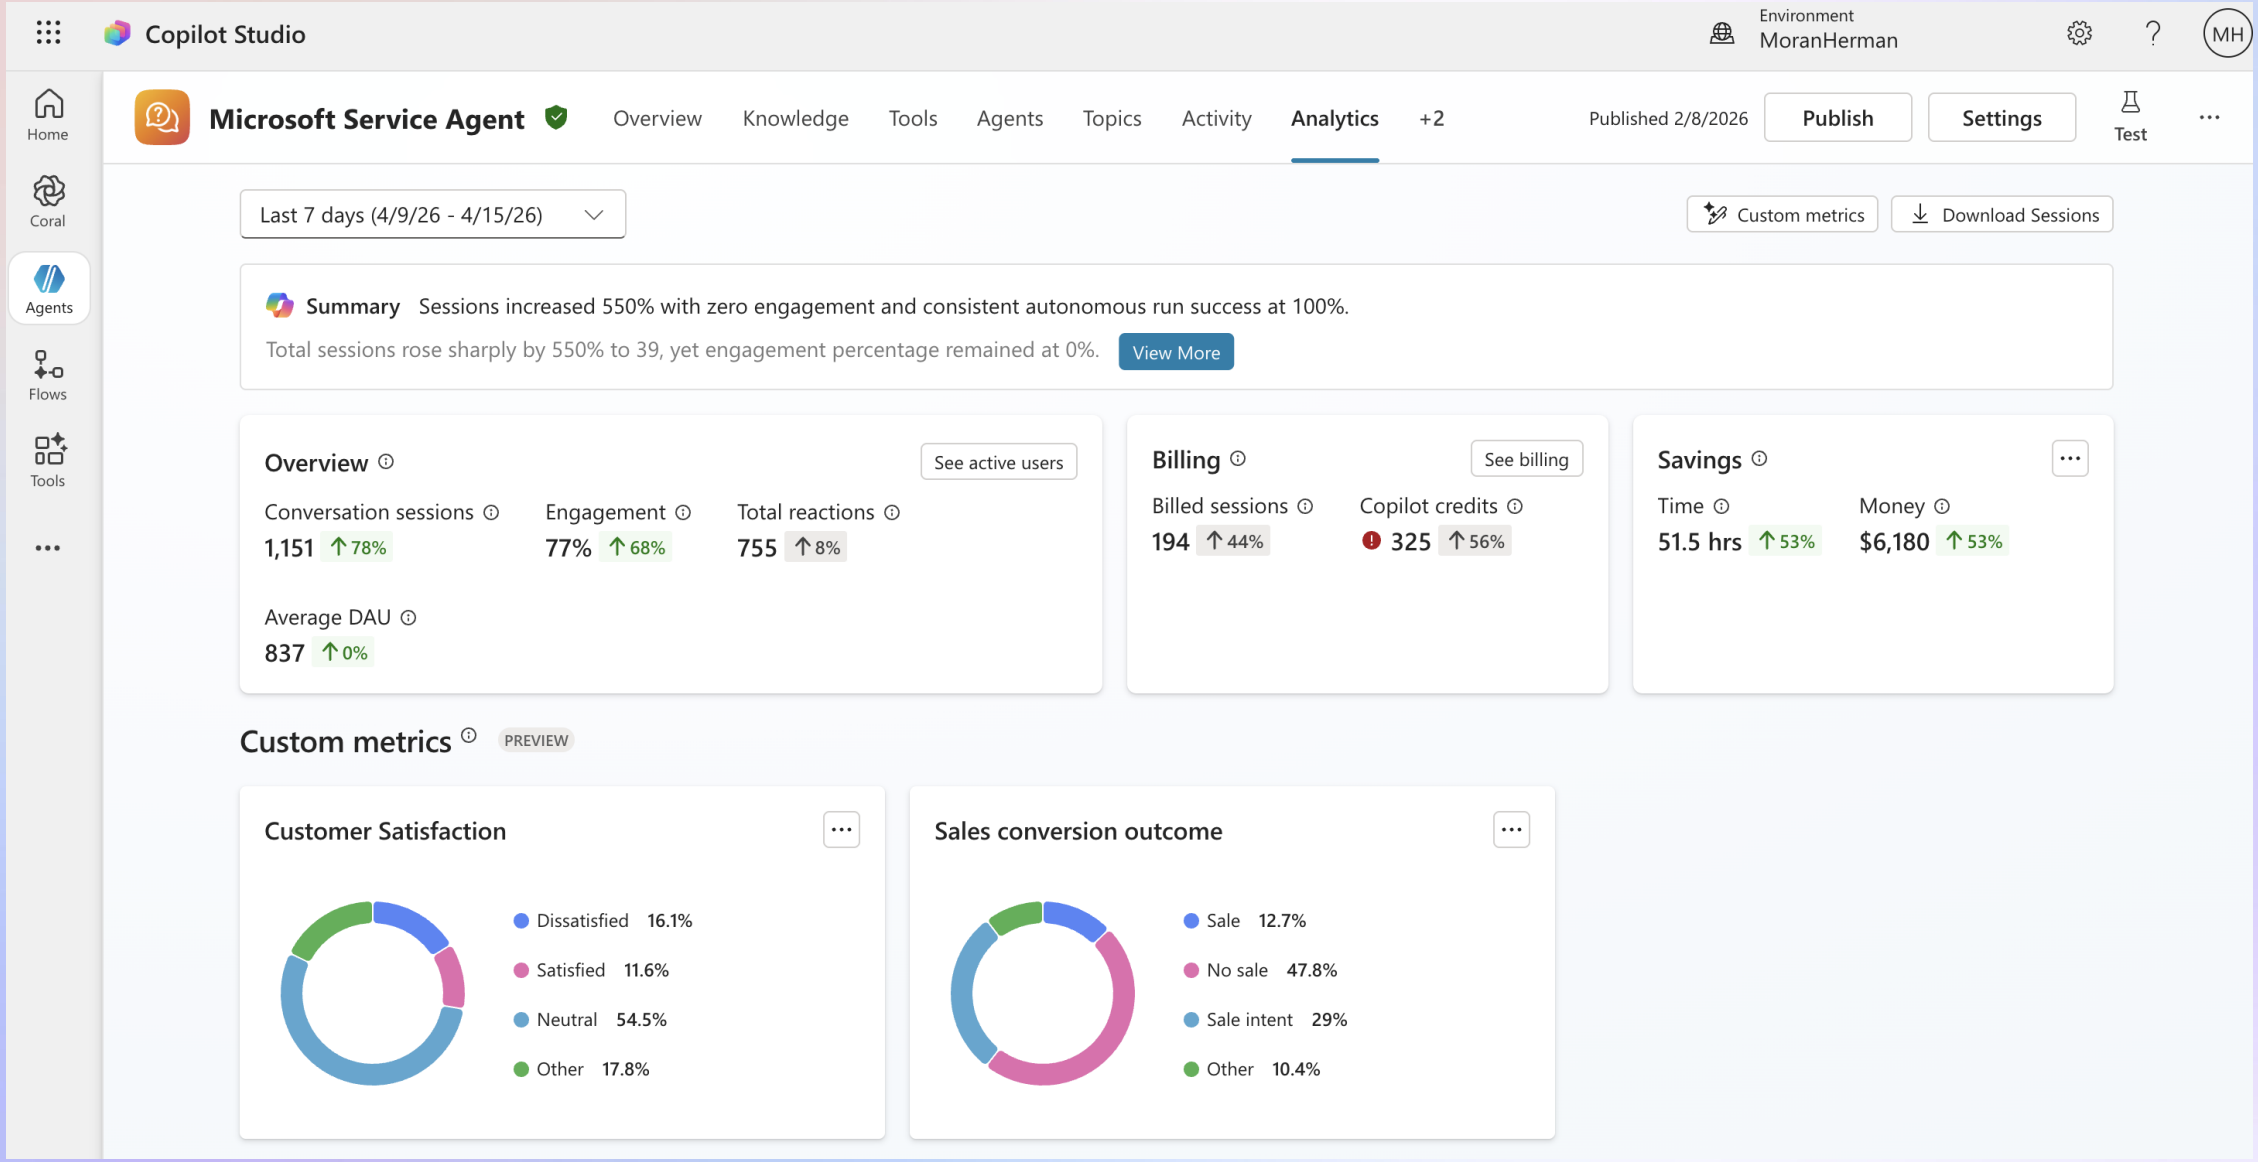

Analytics Viewer Role

Key customer ask

Provide Analysts limited read-only access to the Analytics page.

Analyze agent's performance without the risk to change anything.

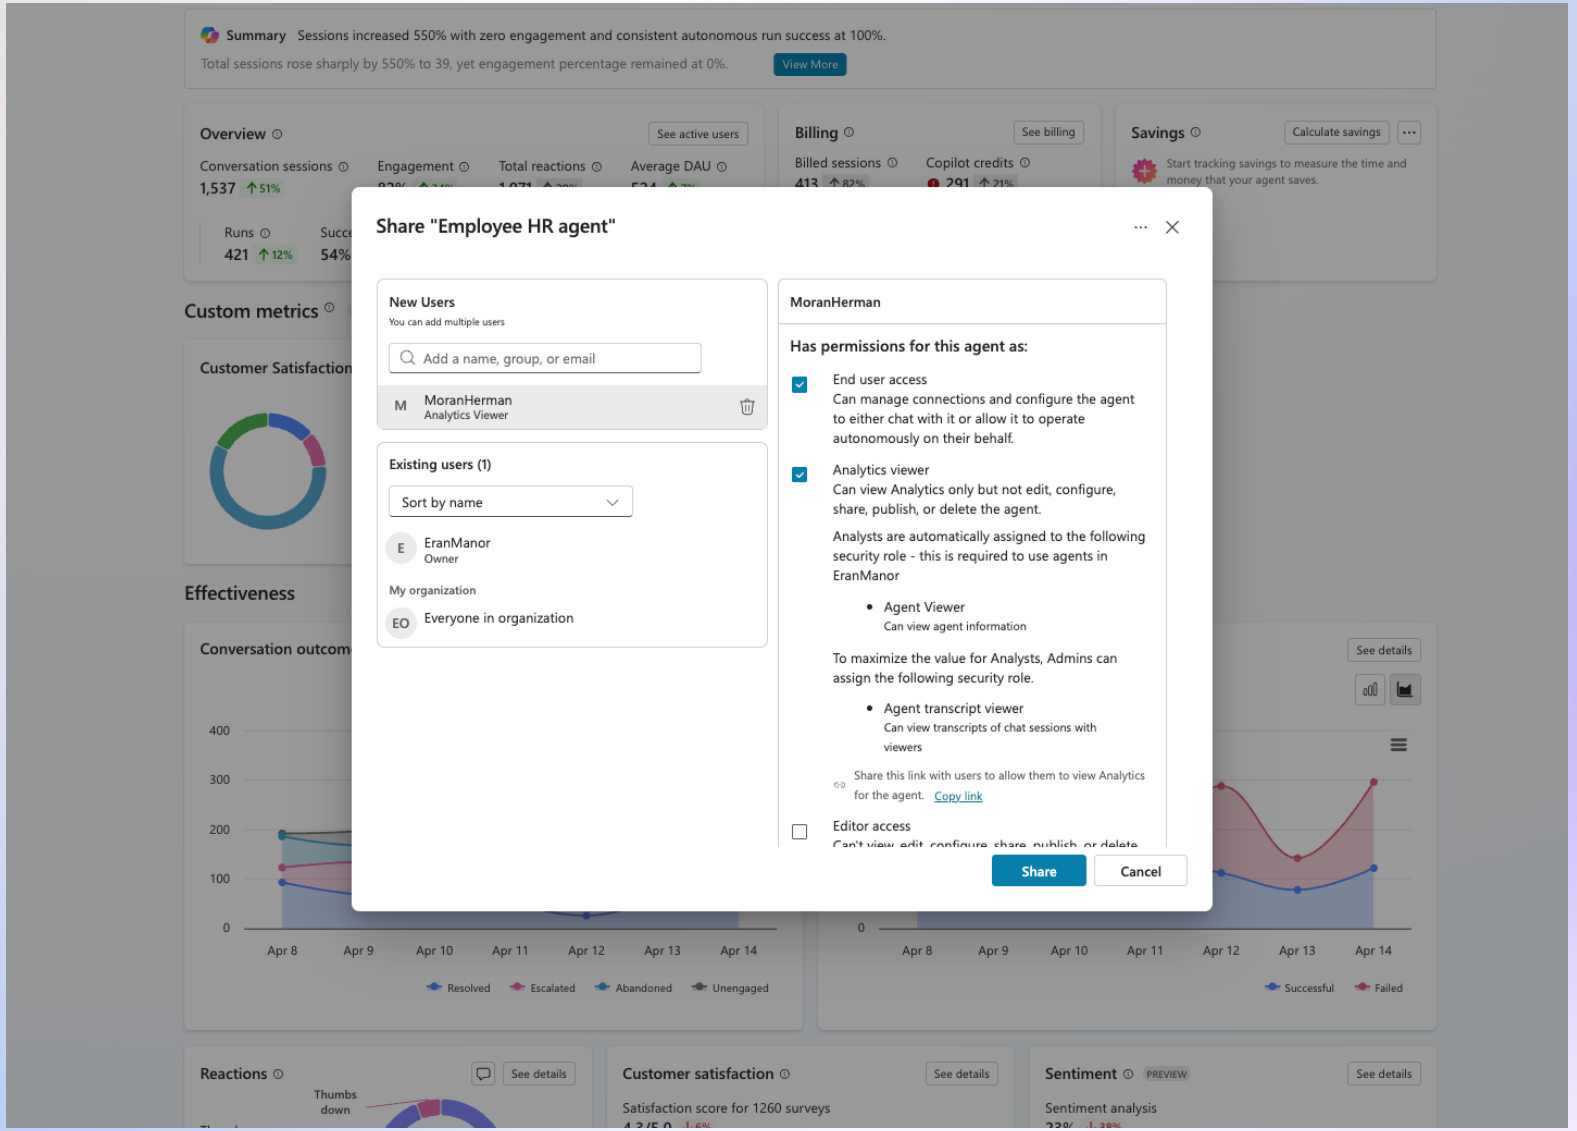

Power CAT Copilot Studio Kit

1. **Agent Value Classification:** AI classifies each agent by type, behavior, and value benefit across 10 defined categories — auto-generates a business value narrative per agent.

2. **Conversation KPIs and Agent Insights Hub:** Aggregated performance data in Dataverse: resolution rates, handoffs, and conversation outcomes — beyond native 30-day analytics.

3. **Agent Inventory:** Tenant-wide view of all agents with usage metrics, feature adoption, auth modes, and exportable data for stakeholder reporting.

4. **Test Automation & Rubrics:** Batch test agents with configurable rubrics ensuring AI grading aligns with business quality standards before deployment.

What it adds beyond native analytics

- Value benefit classification
- Agent Insights
- Conversation KPI trends
- Stakeholder-ready exports
- Cross-agent portfolio view
- Automated quality gates

Built and maintained by the Power CAT team.

Azure App Insights

1. **Custom event tracking:** Log specific business outcomes: form completions, escalation triggers, successful task handoffs
2. **Latency and performance monitoring:** Track response times, identify degradation early, set alerts before users feel impact
3. **KQL query access:** Custom queries against agent telemetry for specific business insights at any granularity
4. **Azure Monitor integration:** Connect to dashboards, alerting pipelines, and operational monitoring across the full Azure estate

When teams typically add App Insights

- The agent handles high-stakes or time-sensitive transactions
- SLA-level performance visibility is needed
- Native analytics do not capture the key business KPIs
- The deployment spans multiple agent types at scale

Value grows in stages

DAY 1 – 90

Immediate Returns

- Query deflection rises
- Time per interaction drops
- Cost-per-case reduction visible
- Agent proves it can handle work

90 DAYS – 1 YEAR

Operational Impact

- Teams organize around agent
- Error rates and rework decline
- Shift toward higher-value work
- Governance becomes trackable

1 YEAR+

Strategic Value

- Decision velocity increases
- New capabilities emerge
- Agent becomes infrastructure
- Returns compound across cases

5 Steps to Build Measurement Discipline

01	BEFORE YOU BUILD **Define what value looks like**	Align on 1–2 outcomes the agent will move. A clear answer now makes the ROI story undeniable later.
02	BEFORE GO-LIVE **Document today's state**	Three numbers—time per task, error rate, cost per interaction—captured before launch, carry the entire measurement story.
03	AT LAUNCH **Turn measurement on from day one**	Enable Analytics, connect App Insights, install the Kit. Every day of missing telemetry is a gap you cannot recover.
04	AT 30 AND 90 DAYS **Connect data to your value pillars**	Compare movement against baseline. Name the pillar explicitly in every review.
05	ONGOING **Make ROI a rhythm, not a report**	Schedule the 90-day review before the agent goes live. Teams that measure consistently discover expansion opportunities.

Now there is a clear answer, at every stage.

Before building · Define what value looks like	During and after · Track what matters	At scale · Let it compound

Let’s build AI agents that prove their worth.